



ECE408/CS483/CSE408

Applied Parallel Programming

Lecture 7:

Convolution and Constant Memory

Exam 1

- Now scheduled for **Tuesday Oct 10th from 7-9pm.**
- The exam will be in-person, on paper, in ECEB.
- Room assignments will be provided.
- Please inform us immediately if you have a legitimate conflict.

Objective

- To learn convolution, an important parallel computation pattern
 - Widely used in signal, image and video processing
 - Foundational to stencil computation used in many science and engineering applications
 - Critical component of Convolutional Neural Networks (CNNs)
- Important GPU technique
 - Taking advantage of cache memories

Convolution

$$f(x) * g(x) = \int_{-\infty}^{\infty} f(\tau) \cdot g(x - \tau) d\tau$$

$$f[x] * g[x] = \sum_{k=-\infty}^{\infty} f[k] \cdot g[x - k]$$

Convolution Applications


- A popular operation that is used in various forms in signal processing, digital recording, image processing, video processing, computer vision, and machine learning.
- Convolution is often performed as a **filter** that transforms the input signal (audio, video, etc) in some context-aware way.
 - Some filters smooth out the signal values so that one can see the big-picture trend
 - Or Gaussian filters to blur images, backgrounds

Convolution Computation

- An array operation where each output data element is a weighted sum of a collection of neighboring input elements
- The weights used in the weighted sum calculation are defined by an input mask array, commonly referred to as the *convolution kernel*

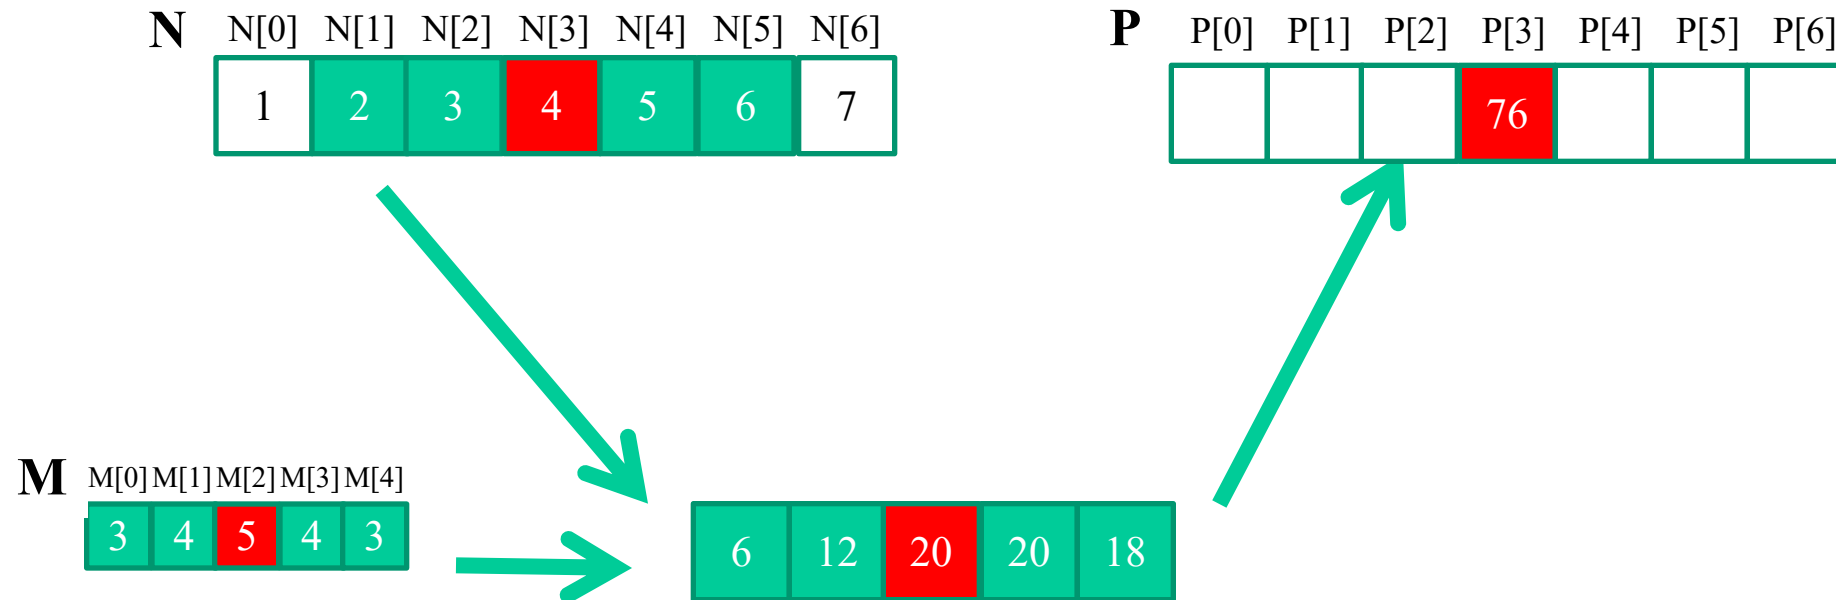
1D Convolution Example

- Commonly used for audio processing
 - **Mask Width** is usually an odd number of elements for symmetry (5 in this example). **Mask Radius** is the number of elements on each side (2 in this example)
- Calculation of P[2]:




1D Convolution Example

- Calculation of $P[3]$



1D Convolution Boundaries

- Calculation of output elements near the boundaries of the input array need to deal with “ghost” elements
 - Different policies (0, replicates of boundary values, etc.)



A 1D Convolution Kernel with Boundary Handling

- This kernel forces all elements outside the valid range to 0
- Each thread calculates one element of P


```
__global__
void convolution_1D_kernel(float *N, float *M, float *P, int Mask_Width, int Width)
{
    int i = blockIdx.x*blockDim.x + threadIdx.x;

    float Pvalue = 0;
    int N_start_point = i - (Mask_Width/2);


    for (int j = 0; j < Mask_Width; j++) {
        if (((N_start_point + j) >= 0) && ((N_start_point + j) < Width)) {
            Pvalue += N[N_start_point + j]*M[j];
        }
    }

    P[i] = Pvalue;
}
```


2D Convolution



2D Convolution Boundary Condition



2D Convolution – Ghost Cells



What does this mask accomplish?

$$\mathbf{M} = \frac{1}{273} \times$$

| | | | | |
|---|----|----|----|---|
| 1 | 4 | 7 | 4 | 1 |
| 4 | 16 | 26 | 16 | 4 |
| 7 | 26 | 41 | 26 | 7 |
| 4 | 16 | 26 | 16 | 4 |
| 1 | 4 | 7 | 4 | 1 |

Assume input N is a grayscale image

What does this mask accomplish?

$\mathbf{M} =$

| | | | | |
|----|----|---|----|----|
| -1 | -1 | 1 | -1 | -1 |
| -1 | -1 | 1 | -1 | -1 |
| -1 | -1 | 1 | -1 | -1 |
| -1 | -1 | 1 | -1 | -1 |
| -1 | -1 | 1 | -1 | -1 |

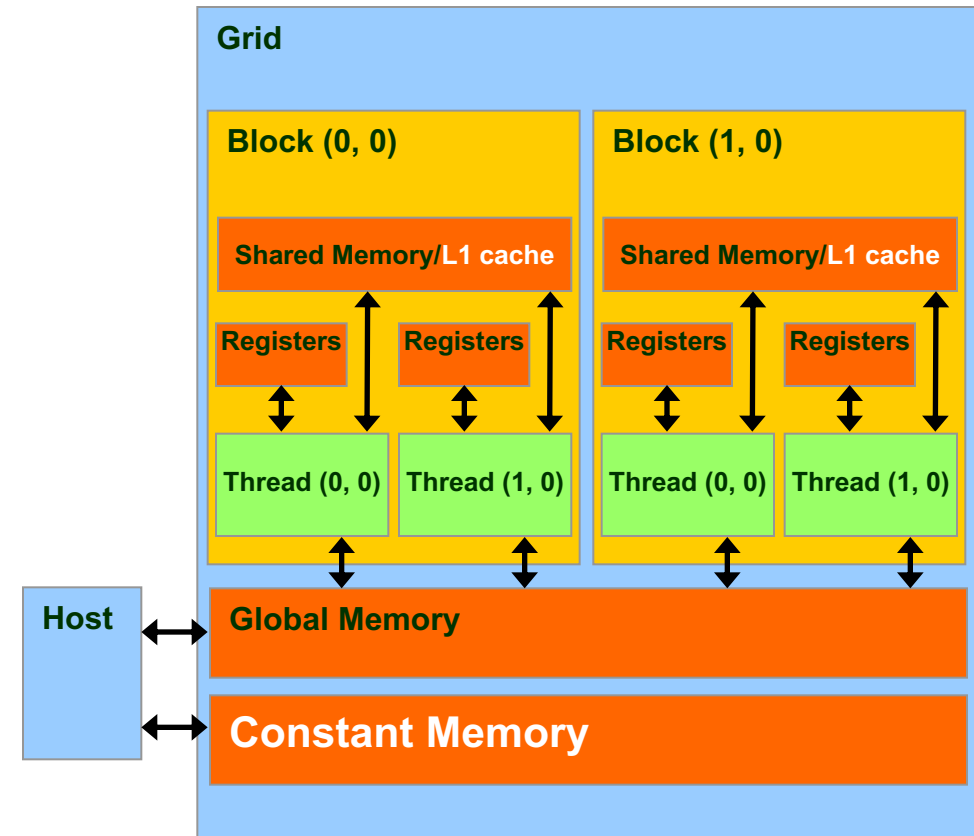
Assume input N is a grayscale image, what does the output P represent?

Access Pattern for M

- Elements of M are called **mask** (kernel, filter) **coefficients**
 - Calculation of all elements of P need M
 - M is not changed during grid execution
- Bonus - M elements are accessed in the same order when calculating all P elements
- M is a good candidate for **Constant Memory**

Programmer View of CUDA Memories

- Each thread can:
 - Read/write per-thread **registers (~1 cycle)**
 - Read/write per-block **shared memory (~5 cycles)**
 - Read/write per-grid **global memory (~500 cycles)**
 - Read/only per-grid **constant memory (~5 cycles with caching)**



Memory Hierarchies

- Review: If all data were in global memory, the execution speed of GPUs would be limited by the global memory bandwidth
- We used shared memory in tiled matrix multiplication to reduce this limitation
- An important part of the solution: caches and constant memory

HW Caches Store Lines of Memory

Recall: memory is optimized for bursts

- contain some number of bit, say **1024 bits (128B)**
- consecutive (linear) addresses.
- Let's call a single burst a **line**.

What's a **cache**?

- An **array of cache lines** (and tags).
- Memory **read produces** a **line**,
- **cache stores** a **copy** of the line, and
- tag records line's memory address.

Memory Accesses Show Locality

An executing program

- loads and store data from memory.
- **Consider sequence of addresses** accessed.

The **Sequence** usually **shows** two types of **locality**:

- **spatial**: accessing **X** implies accessing **X+1** (and X+2, and so forth) **soon**
- **temporal**: accessing **X** implies accessing **X again soon**

Caches improve performance for both types.

Shared Memory vs. Cache

- Caches vs. shared memory
 - Both on-chip, with similar performance
 - As of Nvidia Volta generation, both using the same physical resources, allocated dynamically!

What's the difference?

- Programmer controls shared memory contents (explicit)
- Hardware determines contents of cache (implicit).

GPU Has Both Constant and L1 Caches

To support writes (modification of lines),

- **changes** must be **copied back to memory**, and
- cache must **track** modification **status**.
- **L1 cache** in GPU (for global memory accesses) **supports writes**.

Cache for constant / texture **memory**

- Special case: **lines are read-only**
- Enables higher-throughput access than L1 for common GPU kernel access patterns.

How to Use Constant Memory

Host code is similar to previous versions, but...

Allocate device memory for M (the mask)

- outside of all functions
- using `__constant__`
(tells GPU that caching is safe).

For copying to device memory, use

```
cudaMemcpyToSymbol(dest, src, size, offset = 0, kind =  
cudaMemcpyHostToDevice)
```

- with destination defined as `__constant__`

Host Code Example

```
// global variable, outside any kernel/function
__constant__ float Mc[MASK_WIDTH][MASK_WIDTH];

// Initialize Mask
float Mask[MASK_WIDTH][MASK_WIDTH]
for(unsigned int i = 0; i < MASK_WIDTH * MASK_WIDTH; i++) {
    Mask[i] = (rand() / (float)RAND_MAX);
    if(rand() % 2) Mask[i] = - Mask[i]
}
cudaMemcpyToSymbol(Mc, Mask, MASK_WIDTH*MASK_WIDTH*sizeof(float));

ConvolutionKernel<<<dimGrid, dimBlock>>>(Nd, Pd);
```


Two vertical lines, one blue and one orange, are positioned on the left side of the slide.

**ANY MORE QUESTIONS?
READ CHAPTER 7**